

Vision-Based Reinforcement Learning for Autonomous Drone Racing

Robert Oxley

Autonomous drone racing requires an agent to fly through a sequence of gates within course boundaries using exteroceptive and proprioceptive sensors onboard. Here I present a PPO agent that uses a camera and an inertial measurement unit (IMU) which outputs body rates and continuous thrust at 120 Hz. The policy is trained on privileged ground-truth states, then deployed on noisy monocular perception retaining ~65% of the privileged-state completion rate. This distribution shift is addressed through a perception pipeline combining a learned pose-estimation model, Kalman-filter, Perspective-N-Point, and uncertainty-gated gate association. This autonomy stack successfully completed VQ1 and progressed partway through VQ2 in the Anduril AI Grand Prix, a global autonomous drone racing competition that challenges engineers to prove their autonomy software under realistic, simulated flight conditions.

1 Problem & Approach

Autonomous drone racing is a complex benchmark that can help gauge progress on complex perception, planning, and control algorithms. The drone must perceive, reason, and act upon data within milliseconds, all onboard a computationally limited platform [1]. As such, a conceptual pipeline could involve sensor input, sensor fusion & state estimation, trajectory generation, then a control command is given; leading to the actuation of the drone’s motors.

The Anduril AI Grand Prix featured 2 virtual qualifiers in partnership with the Drone Champions League (DCL): VQ1 consisting of less than 10 gates that had minimal distractions and clear visibility and VQ2 which had less than 20 gates, lighting changes, 3D objects and obstacles. Gates had a 1.5m inner opening and were clear and easily identifiable through the forward-facing first-person camera provided. The simulator provided a rigid-body drone flight model that included thrust generation, aerodynamic drag, gravity, and collision physics with an update rate of 120 Hz. No external localization system was provided nor allowed; all state estimation had to rely on onboard sensing.

In 2025, the ‘MonoRace’ approach [2] won the 2025 A2RL Drone Championship hosted by A2RL and DCL. This system utilized a single rolling-shutter CMOS camera and an Inertial Measurement Unit (IMU) to operate at 500 Hz to emphasize a high-speed, agile flight autonomy stack. The approach I present adopts a similar reward structure to MonoRace.

The approach trains a Proximal Policy Optimization (PPO) algorithm on a 34-dimensional state which outputs collective thrust and body rate commands for an underactuated quadrotor at 120 Hz, under a reward prioritizing reliable gate clearance over lap time. Training runs in a modified Flightmare [3] simulator, a C++ and Python physics simulation and rendering engine for autonomous drones, across a three-phase curriculum of increasing complexity. The drone is trained on privileged ground-truth gate state and then deployed on monocular vision at inference.




Fig. 1. Conceptual Control Architecture for an autonomous drone given the input of an IMU & Camera

2 Results

To evaluate the model, a few metrics are defined:

- Completion Rate: Fraction of runs that crossed gate index N , aggregated over 100 different course variations.
- Residual: $P_{vision} - P_{GT}$, every timestep that the target gate was detected in. P is defined as the gate position and is binned by true distance.
- X/Y and Y/Z trajectories to measure vertical and lateral movements during inference (color-coded with speed).
- X/Z spread of the position where the drone crossed the gate (green for cross, red for hits).

During evaluation, 100 rollouts were performed with 96 course generations unseen by the model during training, the maximum allocated gates were 14, and the maximum time allowed was 40.0s. Fig. 2 shows the completion rates of both ground-truth gate state and inference using vision. Vision tracks ground truth closely early (89% vs 96% at gate 1, 7-point gap) and diverges with depth (52% vs 80% by gate 13, 28-point gap)¹ showcasing that error compounds faster under vision estimates. This demonstrates the policy’s generalization to unseen courses over long-horizon trajectories.

The compounding stems from the distance-dependent perception error that is characterized in the residual graph (Fig. 3). Most notably, the Y-axis has a mean error of -15 when detections are beyond 15 meters, a gate at 15m spans only ~ 18 px in the image, so small corner-localization errors translate into large depth errors, the source of the distance-dependent residual.

Fig. 4 shows a representative run threading the course, colored by speed, while Fig. 5 shows crossings near the opening center.

¹ Gate-14 completion is episode length limited: 21 GT and 18 vision runs reached gate 13 but hit the 40s cap before threading the final gate.




Fig. 2. Completion Rate % per-gate of ground truth (denoted in blue) and vision (denoted in green)




Fig. 3. Vision Residual aggregated over 100 rollouts





Fig. 4. Single trajectory on both the X/Y and Y/Z threading each gate in a course with only vision inference.

Fig. 5. Crossings of a single trajectory near the opening center on only vision inference.



3 Environment

As stated earlier, training and evaluation runs in a modified version of Flightmare. The most notable changes were made to prevent “reward hacking” from the PPO policy. These changes include:

Two-sided gate collision detection to prevent the policy from crossing a gate and then colliding with the object. In training, the negative reward from colliding with a gate is $(-0.5 \cdot n_{hits})$. This is because stopping the episode due to a collision would give PPO, an on-policy algorithm which does not store experience from previous iterations [4], less data to step towards the positive gradient.

Dynamic lateral/vertical walls to prevent the policy from generating a trajectory that uses wide or high/low routes towards the gate. The lowest possible vertical floor is generated by the formula: $\min(z_{lowest} - (2.0 + 2.0 * phase), -1.0)$. The lowest possible lateral wall is generated by: $(\min(x_{lowest}, 0) - random(3.0, 6.0))$ and max is used for the right side. This helps in the event (such as VQ2 in the Anduril AI Grand Prix) where a course that is inside or with obstacles surrounding the gates is present, without losing speed due to obstacle avoidance.

Domain randomization that can be activated through a flag (in this case, during phase 3 of the training process). This updates the drone’s dynamics every episode reset, following the approach of [5], which samples each parameter independently which includes mass, thrust, motor time constant, kappa. This helped to improve sim-to-real transfer and was the primary reason the policy could interact with both the Flightmare simulator and DCL.

The monocular RGB camera is also embedded into the drone, a 20° uptilt, 0.3m z offset, vertical FOV of 90, and outputs images that are 640x360. The environment also includes a flag to configure the initial random rotation of the drone at the starting state, or to launch from the ground. Finally, the entire simulator was updated to Python 3.9, Ubuntu 20.04, and CUDA 13.1 for newer Docker support and Stable-Baselines3/PyTorch support.

4 Model / Training

Model

The policy uses the PPO algorithm from Stable-Baselines3, which is an on-policy, actor-critic reinforcement learning algorithm designed for stable and efficient policy updates. An RL algorithm that is stable is critical for a quadrotor, as they are inherently unstable systems which require an impossible model of each aspect of the drone’s dynamics [6].

The policy and value function’s hidden layers each use a 64x64 MLP while the input is a 34-dimensional observation state and the output is 4 continuous control mappings. The observations are defined as follows:

$$x = [x, y, z, q_R, v_x, v_y, v_z, \omega_p, \omega_q, \omega_r, u_{prev}, (g_x, g_y, g_z)^4]$$
$$u = [t, \phi, \theta, \psi] \in [-1, 1]$$

Here, (x, y, z) equals the position of the gate estimate relative to the drone-body. Q_R is the 9-dimensional rotation matrix which was chosen over quaternions since it is a continuous embedding of SO(3) rather than a discontinuity that degrades network approximation due to the double cover of a quaternion [7].

The observations (v_x, v_y, v_z) and $(\omega_\theta, \omega_\phi, \omega_\psi)$ are the linear and angular velocities of the quadrotor, u_{prev} is the previous action. $(g_x, g_y, g_z)^4$ Are the 4 (x, y, z) coordinates of each corner of the gate estimate that allow the policy to thread a gate without hitting the frame. Finally, the output, $u = [t, \phi, \theta, \psi]$ represents the continuous thrust, roll, pitch, and yaw body rates of the quadrotor.

Each observation is normalized by the running mean/std on the state, but the angular velocities are excluded and normalized by $\omega_{\max}$ instead such that it is able to handle different magnitudes of velocities without being out of distribution.

Training

Training a policy on navigating autonomously through a series of gates is done by running 3 phases of training, which are 1.2×10^8 timesteps for phases 1 and 2, and 3.2×10^8 for phase 3. They are warm started from each phase’s trained weights/normalizer, and are as follows:

Phase 1:

Forward-facing gates, with small lateral/vertical shifts and a time limit of 20.0s.

Phase 2:

Rotated gates with longer lateral/vertical shifts and a time limit of 20.0s. Deviations from the camera result in a dead-reckoned gate estimate to match realistic expectations and only updated every 30Hz to not rely on information each timestep.

Phase 3:

Sharper rotations of gates with steeper lateral/vertical shifts. Deviating from the forward vector of the current gate is penalized more harshly (-0.1 rather than -0.01) and gate crossing bonus is increased from 10 to 30. Structured noise is added to the gate estimates for robustness towards estimates that may be slightly wrong. Additionally, Domain Randomization is introduced which improves sim-2-real transfer, and for the scope of this project sim-2-sim transfer.

Each phase has a constant chance of encountering a gate at the same vertical level as the previous, which in order are [80%, 60%, 40%]. This, along with the vertical and lateral shifts cover most of the edge cases for a gate within the scope of the task and specifications given.

Each phase uses curriculum learning to increment the number of gates during training, from $1 \rightarrow 14$. This is incremented when success clears a threshold formula of $\min(0.65, n^{0.85})$, where n is the amount of gates and 0.85 represents the per-gate reliability that has to be met, maximum is capped at 65% success due to the risk of overfitting a policy for early gates. This represents the horizon length per task difficulty increase, allowing the model to get positive rewards after threading n gates before increasing the length.

Entropy is handled through a reverse entropy schedule. It will start from 0.005 and decay as the gated curriculum advances, towards 0.001.

5 Perception

During training the model received privileged ground-truth gate estimates, however during inference where the stack is restricted to onboard sensors there must exist a pipeline to convert the monocular camera image into the gate estimates.

This autonomy stack utilizes a fine-tuned YOLO pose estimation model to give the 4 inner corner key points of a detected gate in the format $\{tl, tr, br, bl\}$. A confidence threshold of 70% is applied, as a false positive applied to a complex system can be catastrophic, therefore for precision the vision model must emit confidence values above it.

Given the gates known geometry (1.5x1.5 inner width/height), the camera intrinsic and extrinsics, IPPE-Square can be used during Perspective-n-Point (PnP) computation to forego needing a learned depth net or any scale ambiguity [8]. When a detection is not found, the previous detection is dead reckoned until a new detection is found.

A Kalman filter is utilized to update the predictions of a gate estimate rather than raw PnP. IPPE-Square returns two candidate pose solutions, one of which is a flipped version of the frame, therefore raw updates to the gate estimate could result in an unwanted shift in values. By capturing the process noise vector P by incrementing each time a gate is not detected and decreasing when a gate is detected alongside a measurement noise vector R , the gain for an update becomes $K = \frac{P}{P+R}$.

This uncertainty and noise vector is also used for an uncertainty-gated association algorithm, which rejects any detection that diverges from N standard deviation of the uncertainty, which results in the formula $N \cdot \sqrt{P + R}$. The algorithm also rejects any gate that is detected which is similar to the previous gate to prevent a relapse of the gate crossing.

As shown in Fig. 3, the detections of a gate diverge as it approaches 2m of the gate, therefore a gate detection is rejected whose depth is less than 2m and relies on dead reckoning. This is also matched in training such that the transition to vision inference is seamless.

Finally, a stack was made to build a pose-estimation dataset in a YOLO format during data-collection rollouts that are auto-labelled from GT projections, which provided a way to gather data without needing human annotation.




Fig. 6. Conceptual diagram of the perception pipeline during inference.

6 References & Notes

- [1]: D. Hanover *et al.*, “Autonomous Drone Racing: A Survey,” *IEEE transactions on robotics*, pp. 1–24, Jan. 2024, doi: 10.1109/tro.2024.3400838.
- [2]: S. A. Bahnam *et al.*, “MonoRace: Winning Champion-Level Drone Racing with Robust Monocular AI,” *arXiv.org*, 2026. <https://arxiv.org/abs/2601.15222>
- [3]: Y. Song, S. Naji, E. Kaufmann, A. Loquercio, and D. Scaramuzza, “Flightmare: A Flexible Quadrotor Simulator,” *arXiv.org*, 2020. <https://arxiv.org/abs/2009.00563>
- [4]: J. Schulman, F. Wolski, P. Dhariwal, A. Radford, and O. Klimov, “Proximal Policy Optimization Algorithms,” *arXiv.org*, Aug. 28, 2017. <https://arxiv.org/abs/1707.06347>
- [5]: R. Ferede, B. Till, E. Lucassen, De W. Christophe, and de C. G. C. H. E, “One Net to Rule Them All: Domain Randomization in Quadcopter Racing Across Different Platforms,” *arXiv.org*, 2025. <https://arxiv.org/abs/2504.21586>
- [6]: G. Cano Lopes, M. Ferreira, A. da Silva Simões and E. Luna Colombini, "Intelligent Control of a Quadrotor with Proximal Policy Optimization Reinforcement Learning," pp. 503-508, 2018, doi: 10.1109/LARS/SBR/WRE.2018.00094.
- [7]: Y. Zhou, C. Barnes, J. Lu, J. Yang, and H. Li, “On the Continuity of Rotation Representations in Neural Networks,” *arXiv.org*, 2018. <https://arxiv.org/abs/1812.07035>
- [8]: T. Collins and A. Bartoli, “Infinitesimal Plane-Based Pose Estimation,” *International Journal of Computer Vision*, vol. 109, no. 3, pp. 252–286, July 2014, doi: 10.1007/s11263-014-0725-5